

Optimizers in deep learning - Practical study

CPE 727 - Deep Learning

Ana Clara Loureiro Cruz Bruno Coelho Martins Emre Aslan Felipe

Barreto Andrade (anaclaralcruz@poli.ufrj.br)

bruno.martins@smt.ufrj.br

emre@gta.ufrj.br

felipebarretoandrade@poli.ufrj.br)

November 10, 2025


Table of Contents

1 Gradient Descent

- ▶ Gradient Descent
- ▶ Adam and Its Variants
- ▶ Learning-rate (LR) Schedulers
- ▶ Limited BFGS [1]
- ▶ Comparing different types of optimizers
- ▶ References

Gradient Descent (GD) Optimization

1 Gradient Descent



[GitHub Link](#)

Gradient Descent


1 Gradient Descent

Table: Comparison of Hyperparameters and Key Differences between Gradient Descent (GD) and Stochastic Gradient Descent (SGD)

Aspect	Gradient Descent (GD)	Stochastic Gradient Descent (SGD)
Learning Rate	0.1 (fixed)	0.1 (fixed)
Number of Epochs	20	20
Batch Size	500 (full dataset)	1 (single sample)
Gradient Computation	Full dataset (500 samples)	Single sample per iteration
Updates per Epoch	1 update	500 updates
Computational Cost per Update	High (processes all samples)	Low (processes one sample)
Convergence Behavior	Smoother, more stable	Noisier, faster per iteration
Implementation Detail	Uses entire dataset for gradient	Randomly samples one data point

Standard Sgd vs Mini-batch Sgd

1 Gradient Descent



Standard Sgd vs Mini-batch Sgd

1 Gradient Descent


[GitHub Link](#)

Aspect	Standard SGD	Mini-batch SGD
Learning Rate	0.1 (fixed)	0.1 (fixed)
Number of Epochs	20	20
Batch Size	1 (single sample)	32 (multiple samples)
Gradient Computation	Single sample	32 samples per iteration
Updates per Epoch	500 updates	16 updates (500/32)
Computational Cost	Low (one sample per update)	Moderate (32 samples per update)
Convergence Behavior	Noisier, high variance	Smoother, reduced variance
Implementation Detail	Randomly samples one data point	Randomly samples 32 data points

Table: Comparison of Hyperparameters and Key Differences between Standard SGD and Mini-batch SGD

SGD with Momentum

1 Gradient Descent



[GitHub Link](#)

SGD with Momentum

1 Gradient Descent

Aspect	Standard SGD	SGD with Momentum
Learning Rate	0.1 (fixed)	0.1 (fixed)
Number of Epochs	20	20
Batch Size	1 (single sample)	1 (single sample)
Momentum Coefficient	None	0.9
Gradient Update	Direct gradient: $\theta \leftarrow \theta - \eta \nabla J$	Velocity-based: $v \leftarrow \mu v - \eta \nabla J$, $\theta \leftarrow \theta + v$
Updates per Epoch	500 updates	500 updates
Convergence Behavior	Noisier, high variance	Smoother, faster due to momentum
Implementation Detail	Updates with raw gradient	Uses velocity to accelerate gradients

Table: Comparison of Hyperparameters and Key Differences between Standard SGD and SGD with Momentum

Table of Contents

2 Adam and Its Variants

- ▶ Gradient Descent
- ▶ Adam and Its Variants
- ▶ Learning-rate (LR) Schedulers
- ▶ Limited BFGS [1]
- ▶ Comparing different types of optimizers
- ▶ References

Experiment Setup: Rosenbrock Function

2 Adam and Its Variants

- All experiments are conducted on the **Rosenbrock function**, a classic benchmark for testing optimizers.


$$f(x, y) = (1 - x)^2 + 100(y - x^2)^2$$

- This function forms a narrow, curved valley that makes optimization difficult:
 - Gradients are small along the valley but steep across it.
 - Optimizers must balance stability and adaptivity to reach the minimum efficiently.
- It is ideal to visualize how each algorithm handles curvature, adaptivity, and noise.

Effect of Hyperparameters in Adam

2 Adam and Its Variants

Rosenbrock function: $a = 1, b = 100$



Adam hyperparameters: effect of beta2

2 Adam and Its Variants

```
sweep_specs = [
    ("Adam beta2=0.99", {"betas": (0.9, 0.99)}),
    ("Adam beta2=0.95", {"betas": (0.9, 0.95)}),
]
for label, kw in sweep_specs:
    _, losses, _ = run_optimizer(torch.optim.Adam, steps=800, lr=3e-3, **kw)
    plt.semilogy(losses, label=label)
plt.title("Adam hyperparameters: effect of beta2")
```


Effect of Hyperparameters in Adam

2 Adam and Its Variants

- The parameter β_2 controls the smoothing of the variance (second moment):
 - Higher values (0.99) produce smoother but slower updates.
 - Lower values (0.95) react faster to gradient changes.
- Small changes in β_2 can strongly influence training speed and stability.

Effect of Hyperparameters in Adam

2 Adam and Its Variants



Adam vs AdamW (decoupled weight decay)

2 Adam and Its Variants

```
_, loss_adam, _ = run_optimizer(torch.optim.Adam, lr=3e-3, weight_decay=1e-2)
_, loss_adamw, _ = run_optimizer(torch.optim.AdamW, lr=3e-3, weight_decay=1e-2)

plt.semilogy(loss_adam, label="Adam (wd=0.01)")
plt.semilogy(loss_adamw, label="AdamW (wd=0.01)")
plt.title("Coupled L2 (Adam) vs Decoupled (AdamW)")
```

Conceptual Difference: Adam vs AdamW

2 Adam and Its Variants

- Original Adam couples L2 regularization with adaptive updates:

$$\theta_{t+1} = \theta_t - \eta \left(\frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon} + \lambda \theta_t \right)$$


- AdamW decouples weight decay:

$$\theta_{t+1} = (1 - \eta \lambda) \theta_t - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

- This avoids the interference between regularization and learning rates.
- Effect: cleaner weight decay, better generalization (especially in Transformers).

Adam vs AdamW: Convergence Comparison

2 Adam and Its Variants



Loss vs iterations for Adam-family optimizers

2 Adam and Its Variants

```
variants = [
    ("Adam",      torch.optim.Adam,   {"amsgrad": False}),
    ("AMSGrad",   torch.optim.Adam,   {"amsgrad": True}),
    ("RAdam",     torch.optim.RAdam,  {}),
    ("AdamW",     torch.optim.AdamW,  {}),
]
for name, ctor, kw in variants:
    _, losses, _ = run_optimizer(ctor, lr=3e-3, steps=800, **kw)
    plt.semilogy(losses, label=name)
plt.title("Loss vs iterations (log scale)")
```


Comparison Across Adam-family Variants

2 Adam and Its Variants

- The plot shows training loss per iteration (log scale).
- **Adam**: fast early convergence, but can oscillate.
- **AMSGrad**: uses $\max(v_t)$ to ensure monotonic variance estimates.
- **RAdam**: introduces variance rectification to fix the warmup problem.
- **AdamW**: stable convergence and cleaner weight decay.

Comparison Across Adam-family Variants

2 Adam and Its Variants



Step size evolution

2 Adam and Its Variants

```
for name in results:  
    y = results[name] ["dtheta"]  
    plt.plot(y, label=name)  
plt.yscale("symlog", linthresh=1e-5)  
plt.title("Step size over time (symlog)")  
plt.legend()
```


Evolution of Step Size

2 Adam and Its Variants

- $\|\Delta\theta_t\|$ measures how much parameters change at each iteration.
- Large spikes at the beginning reflect **adaptive warmup**.
- **RAdam** shows the highest variance initially.
- **AdamW** and **AMSGrad** stabilize faster and keep smaller steps.
- The **symlog** scale allows both small and large step magnitudes to be visible.

Evolution of Step Size

2 Adam and Its Variants



Parameter-space trajectories

2 Adam and Its Variants

```
X, Y, Z = rosenbrock_grid()
for name in results:
    T = results[name]["traj"]
    plt.plot(T[:,0], T[:,1], label=name)
plt.contour(X, Y, Z, levels=np.logspace(-1,3,24))
plt.title("Rosenbrock: parameter-space trajectories")
```

Parameter-space Trajectories

2 Adam and Its Variants

- The Rosenbrock function forms a curved valley leading to the global minimum.
- **Adam** (blue): follows a relatively smooth trajectory, with minor lateral oscillations.
- **AMSGrad** (purple): exhibits shorter, more consistent steps by using the maximum historical variance to stabilize updates.
- **RAdam** (yellow): shows wider initial swings since the rectification warmup is still adapting.
- **AdamW** (green): maintains a direct and stable path, benefiting from its decoupled weight decay.

Parameter-space Trajectories

2 Adam and Its Variants

Rosenbrock: parameter-space trajectories (zoomed)




Table of Contents

3 Learning-rate (LR) Schedulers

- ▶ Gradient Descent
- ▶ Adam and Its Variants
- ▶ Learning-rate (LR) Schedulers
- ▶ Limited BFGS [1]
- ▶ Comparing different types of optimizers
- ▶ References

Scheduler class

3 Learning-rate (LR) Schedulers

```
class Scheduler:  
    """  
    Scheduler: create PyTorch LR schedulers from a name + params.  
  
    Supported schedulers:  
    LambdaLR, MultiplicativeLR, StepLR, MultiStepLR, ConstantLR, LinearLR,  
    ExponentialLR, PolynomialLR (builtin if available, else Lambda fallback),  
    CosineAnnealingLR, ChainedScheduler, SequentialLR, ReduceLROnPlateau,  
    CyclicLR, OneCycleLR.  
  
    The class returns the scheduler object already constructed and ready to be  
    stepped in the training loop.  
    """
```

Experiment Overview

3 Learning-rate (LR) Schedulers

This experiment trains a simple MLP model on the Breast Cancer dataset with preprocessing and supports flexible learning rate schedulers (similar to Breast Cancer MLP Experiment).

Experiment - Default Parameters

3 Learning-rate (LR) Schedulers

- epochs: 100
- batch_size: 32
- learning_rate: 0.05
- hidden_size: 64
- feature_strategy: onehot
- target_strategy: binary
- handle_missing: drop
- device: cpu
- scheduler_name: CosineAnnealingLR
- scheduler_params: {} (JSON string)

Experiment - Output

3 Learning-rate (LR) Schedulers

- Train loss per epoch plot: train_loss_per_epoch_<SCHEDULER_NAME>.png
- Learning rate per epoch plot: lr_per_epoch_<SCHEDULER_NAME>.png

Running the Experiment

3 Learning-rate (LR) Schedulers

Run the experiment using Python module syntax and the CLI script:

```
python -m src.experiments.LRSchedulerExperiment.lrscheduler_experiment.cli \
--scheduler <SCHEDULER_NAME> \
--scheduler-params '<JSON_PARAMS>'
```



LambdaLR — Flexible functional schedules

3 Learning-rate (LR) Schedulers

- **lr_lambda (callable)** — `function(epoch)` that returns a multiplicative factor.
 - Default: `lambda epoch: 1/(1+0.1*epoch)`
 - Effect: completely custom per-epoch scaling (useful for bespoke decays).

LambdaLR - Curves

3 Learning-rate (LR) Schedulers



- Parameters used: $\text{lr_lambda} = 1/(1 + 0.1 * \text{epoch})$ (default)
- Final Train Loss: 0.0329
- Final Test Loss: 1.4957
- Final Test Accuracy: 71.43%



MultiplicativeLR — Repeated multiplicative updates

3 Learning-rate (LR) Schedulers

- **factor (float)** — multiplicative factor applied each epoch.
 - Default: 0.95
 - Effect: $\text{lr}_{t+1} = \text{lr}_t \times \text{factor}$ (exponential-like decay).
- **lr_lambda (callable)** — alternative callable (defaults to constant factor).

MultiplicativeLR - Curves

3 Learning-rate (LR) Schedulers



- Parameters used: factor = 0.95 (default)
- Final Train Loss: 0.0337
- Final Test Loss: 1.2416
- Final Test Accuracy: 71.43%



StepLR — Discrete drops

3 Learning-rate (LR) Schedulers

- **step_size (int)** — epochs between drops.
 - Default: 30
- **gamma (float)** — multiplicative drop factor.
 - Default: 0.1

StepLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: `step_size = 33`, `gamma = 0.3`
- Final Train Loss: 0.0344
- Final Test Loss: 1.7086
- Final Test Accuracy: 67.86%



MultiStepLR — Custom step milestones

3 Learning-rate (LR) Schedulers

- **milestones (list[int])** — epochs where LR is reduced.
 - Default: [30, 60, 90]
- **gamma (float)** — multiplicative factor at each milestone.
 - Default: 0.1

MultiStepLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: milestones = [30, 80], gamma = 0.3
- Final Train Loss: 0.0279
- Final Test Loss: 1.1761
- Final Test Accuracy: 73.21%



ConstantLR — Constant schedule (warmup-like)

3 Learning-rate (LR) Schedulers

- **factor (float)** — multiplier applied during the constant period.
 - Default: 0.1
- **total_iters (int)** — number of iterations the factor is kept.
 - Default: `max(1, int(0.05 * total_steps))` (5% of total steps if known)

ConstantLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: factor = 0.1, total_iters = 10
- Final Train Loss: 0.0281
- Final Test Loss: 1.7522
- Final Test Accuracy: 69.64%



LinearLR — Linear warmup

3 Learning-rate (LR) Schedulers

- **start_factor (float)** — starting multiplier of base LR.
 - Default: 0.1
- **total_iters (int)** — warmup duration in iterations.
 - Default: `max(1, int(0.05 * total_steps))`

LinearLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: `start_factor = 0.1, total_iters = 10`
- Final Train Loss: 0.0264
- Final Test Loss: 1.6642
- Final Test Accuracy: 82.14%



ExponentialLR — Exponential decay

3 Learning-rate (LR) Schedulers

- **gamma (float)** — multiplicative decay factor per epoch/step.
 - Default: 0.95

ExponentialLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: gamma = 0.8
- Final Train Loss: 0.2508
- Final Test Loss: 0.6361
- Final Test Accuracy: 73.21%



PolynomialLR — Polynomial decay

3 Learning-rate (LR) Schedulers

- **power (float)** — exponent of the polynomial.
 - Default: 2.0
- **total_iters (int)** — total number of iterations the decay spans.
 - Default: total_steps (if known), else number of epochs
- Effect: $\text{lr}_t = \text{lr}_0 \times (1 - t/\text{max_iter})^{\text{power}}$

PolynomialLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: power = 2.0, max_iter = 200
- Final Train Loss: 0.0312
- Final Test Loss: 1.3706
- Final Test Accuracy: 76.79%



CosineAnnealingLR — Smooth cosine annealing

3 Learning-rate (LR) Schedulers

- **T_max (int)** — number of epochs or steps in one cycle.
 - Default: num_epochs
- **eta_min (float)** — minimum learning rate (floor).
 - Default: 0.0

CosineAnnealingLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: $T_{\text{max}} = 100$, $\eta_{\text{min}} = 1e-6$
- Final Train Loss: 0.0317
- Final Test Loss: 1.9866
- Final Test Accuracy: 62.50%



ChainedScheduler — Combining schedules

3 Learning-rate (LR) Schedulers

- Chains multiple schedulers sequentially (each scheduler runs for its configured duration).
- default factory: **LinearLR** warmup then **CosineAnnealingLR** main.
- No user params required for the default chain (factory builds sensible warmup length).

ChainedScheduler — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: default (Linear warmup + Cosine main)
- Final Train Loss: 0.0267
- Final Test Loss: 1.7336
- Final Test Accuracy: 71.43%



SequentialLR — Sequential composition

3 Learning-rate (LR) Schedulers

- **schedulers**: list of schedulers to run in sequence.
- **milestones**: list of integers indicating when to switch to the next scheduler.
 - Default: [warmup_iters] where warmup_iters 5% of total steps
- Use-case: warmup (LinearLR) → main (CosineAnnealingLR).

SequentialLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: default (LinearLR warmup + CosineAnnealingLR)
- Final Train Loss: 0.0276
- Final Test Loss: 1.7895
- Final Test Accuracy: 71.43%



ReduceLROnPlateau — Metric-driven reductions

3 Learning-rate (LR) Schedulers

- **mode** — “min” or “max”; which direction is “better”.
 - Default: ‘min’ (monitor metrics like validation loss)
- **factor (float)** — LR multiplication factor when reducing.
 - Default: 0.1
- **patience (int)** — epochs without improvement before reducing.
 - Default: 5
- **threshold (float)** — minimal change to count as improvement.
 - Default: 1e-4
- **cooldown (int)** — epochs to wait after reduction.
 - Default: 0
- **min_lr (float)** — lower bound for LR.
 - Default: 0.0
- **eps (float)** — minimal decay to avoid tiny updates.
 - Default: 1e-8

ReduceLROnPlateau — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: mode = min, factor = 0.5, patience = 3
- Final Train Loss: 0.1935
- Final Test Loss: 0.7553
- Final Test Accuracy: 66.07%



CyclicLR — Cyclical policies

3 Learning-rate (LR) Schedulers

- **base_lr** — lower bound of cycle.
 - Default: $0.1 \times \text{initial lr}$
- **max_lr** — upper bound of cycle.
 - Default: $10 \times \text{initial lr}$
- **step_size_up (int)** — iterations to increase from base to max.
 - Default: `max(1, floor(steps_per_epoch/2))`
- **mode** — ‘triangular’, ‘triangular2’, ‘exp_range’.
 - Default: ‘triangular’
- **cycle_momentum** — whether to cycle momentum as well.
 - Default: False

CyclicLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: `base_lr = 0.0001`, `max_lr = 0.001`, `step_size_up = 100`
- Final Train Loss: 0.3535
- Final Test Loss: 0.6226
- Final Test Accuracy: 69.64%



OneCycleLR — Single-cycle super-convergence

3 Learning-rate (LR) Schedulers

- **max_lr (float)** — peak learning rate.
 - Default: $10 \times$ initial lr
- **total_steps (int)** — total number of optimizer steps (batches).
 - Default: $\text{num_epochs} \times \text{steps_per_epoch}$ (must be correct for batch stepping)
- **pct_start (float)** — fraction of total steps spent increasing to max_lr.
 - Default: 0.3
- **anneal_strategy** — ‘cos’ or ‘linear’.
 - Default: ‘cos’
- **div_factor** — initial LR = $\text{max_lr} / \text{div_factor}$.
 - Default: 25.0
- **final_div_factor** — min LR = initial LR / final_div_factor.
 - Default: 1e4
- Note: **step this scheduler every batch**, not per epoch.

OneCycleLR — Curves

3 Learning-rate (LR) Schedulers



- Parameters used: `max_lr = 0.01, total_steps = 3200, batch_size = 32`
- Final Train Loss: 0.0544
- Final Test Loss: 0.7502
- Final Test Accuracy: 69.64%

Table of Contents

4 Limited BFGS [1]

- ▶ Gradient Descent
- ▶ Adam and Its Variants
- ▶ Learning-rate (LR) Schedulers
- ▶ Limited BFGS [1]
- ▶ Comparing different types of optimizers
- ▶ References

Experiment: *VariateHistorySizeInLBFGSExperiment*

4 Limited BFGS [1]

[Link to GitHub](#)

- Goal: isolates the effect of m (memory size) in L-BFGS.
- Dataset: MNIST [2]
- Base model: simple MLP (Multi-Layer Perceptron).
- Evaluation:
 - Average loss.
 - Accuracy.

Experiment: *VariateHistorySizeInLBFGSExperiment*

4 Limited BFGS [1]

[Link to GitHub](#)

The comparison is done as following:

```
for step in range(outer_steps):
    loss = optim.step(closure)
    # Log da perda (como número Python)
    outer_losses.append(float(loss.detach().cpu()))
```

Results

4 Limited BFGS [1]




Table of Contents

5 Comparing different types of optimizers

- ▶ Gradient Descent
- ▶ Adam and Its Variants
- ▶ Learning-rate (LR) Schedulers
- ▶ Limited BFGS [1]
- ▶ Comparing different types of optimizers
- ▶ References

Dataloader: CIFAR-10 [3]

5 Comparing different types of optimizers

[Link to GitHub](#)

```
class LabelNoiseDataset(torch.utils.data.Dataset):
    """Aplica ruido de rótulo com prob p_noise."""
    def __init__(self, base_ds, p_noise: float, num_classes: int = 10, seed: int = 0):
        self.base = base_ds
        self.p = float(p_noise)
        self.C = int(num_classes)
        rng = np.random.RandomState(seed)
        self.flip_mask = rng.rand(len(self.base)) < self.p
        self.rand_labels = rng.randint(0, self.C, size=len(self.base))
```

```
class TestNoiseWrapper(torch.utils.data.Dataset):
    """
    Adiciona ruido gaussiano em pixel space, faz clamp [0,1] e depois normaliza com mean/std.
    """
    def __init__(self, base_ds, sigma: float, mean_=mean, std_=std):
        self.base, self.sigma = base_ds, float(sigma)
        self.normalize = T.Normalize(mean_, std_)
```

Experiment: Optimizers Robustness To Noise Experiment

5 Comparing different types of optimizers

[Link to GitHub](#)

- Goal: compare robustness to noise and memory use among different optimizers in CIFAR-10.
- Base model: ResNet-18 (10 classes).
- Types of noise:
 - **Label noise** in training (with $p = 0.2$).
 - **Input noise** in testing ($\sigma \in \{0, 0.05, 0.1, 0.2\}$).
- Evaluation:
 - Área unther the curve (AUC): model's ability to rank positive examples higher than negative ones.
 - Pick of memory usage CUDA.

Configuration

5 Comparing different types of optimizers

[Link to GitHub](#)

```
def make_optimizer(name, params):
    name = name.lower()
    if name == 'sgd':
        return torch.optim.SGD(params, lr=0.1, momentum=0.0, weight_decay=5e-4)
    if name == 'nesterov':
        return torch.optim.SGD(params, lr=0.1, momentum=0.9, nesterov=True, weight_decay=5e-4)
    if name == 'rmsprop':
        return torch.optim.RMSprop(params, lr=1e-3, alpha=0.99, eps=1e-8, weight_decay=1e-5)
    if name == 'adamw':
        return torch.optim.AdamW(params, lr=3e-4, betas=(0.9,0.999), eps=1e-8, weight_decay=0.01)
    if name == 'radam':
        return torch.optim.RAdam(params, lr=3e-4, betas=(0.9,0.999), eps=1e-8, weight_decay=1e-4)
    if name == 'l-bfgs':
        return torch.optim.LBFGS(params, lr=1.0, history_size=10, line_search_fn=None)
```

Evaluation

5 Comparing different types of optimizers


[Link to GitHub](#)

```
def auc_robustez(acc_by_sigma: dict):
    xs = sorted(acc_by_sigma.keys())
    ys = [acc_by_sigma[x]['acc'] for x in xs]
    area = 0.0
    for i in range(len(xs)-1):
        dx = xs[i+1] - xs[i]
        area += 0.5 * (ys[i] + ys[i+1]) * dx
    return area
```

Results

5 Comparing different types of optimizers


Starting experiment on cuda | EPOCHS=5 | BATCH_SIZE=128 | SUBSET_TRAIN=10000
Optimizers: ['SGD', 'Nesterov', 'RMSProp', 'AdamW', 'RAdam', 'L-BFGS']



Results

5 Comparing different types of optimizers

```
Starting experiment on cuda | EPOCHS=10 | BATCH_SIZE=128 | SUBSET_TRAIN=10000
Optimizers: ['SGD', 'Nesterov', 'RMSProp', 'AdamW', 'RAdam', 'L-BFGS']
```



Results

5 Comparing different types of optimizers

Starting experiment on cuda | EPOCHS=10 | BATCH_SIZE=128 | SUBSET_TRAIN=20000
Optimizers: ['SGD', 'Nesterov', 'RMSProp', 'AdamW', 'RAdam', 'L-BFGS']




Table of Contents

6 References

- ▶ Gradient Descent
- ▶ Adam and Its Variants
- ▶ Learning-rate (LR) Schedulers
- ▶ Limited BFGS [1]
- ▶ Comparing different types of optimizers
- ▶ References

Bibliography

6 References

- [1] PyTorch Core Team, *LBFGS - PyTorch Documentation*, 2025.
Accessed: October 2025.
- [2] Y. LeCun, C. Cortes, and C. Burges, "Mnist handwritten digit database," *ATT Labs [Online]*. Available: <http://yann.lecun.com/exdb/mnist>, vol. 2, 2010.
- [3] A. Krizhevsky, "Learning multiple layers of features from tiny images," tech. rep., 2009.

Optimizers in deep learning - Practical study

*Obrigado pela Atenção!
Alguma Pergunta?*